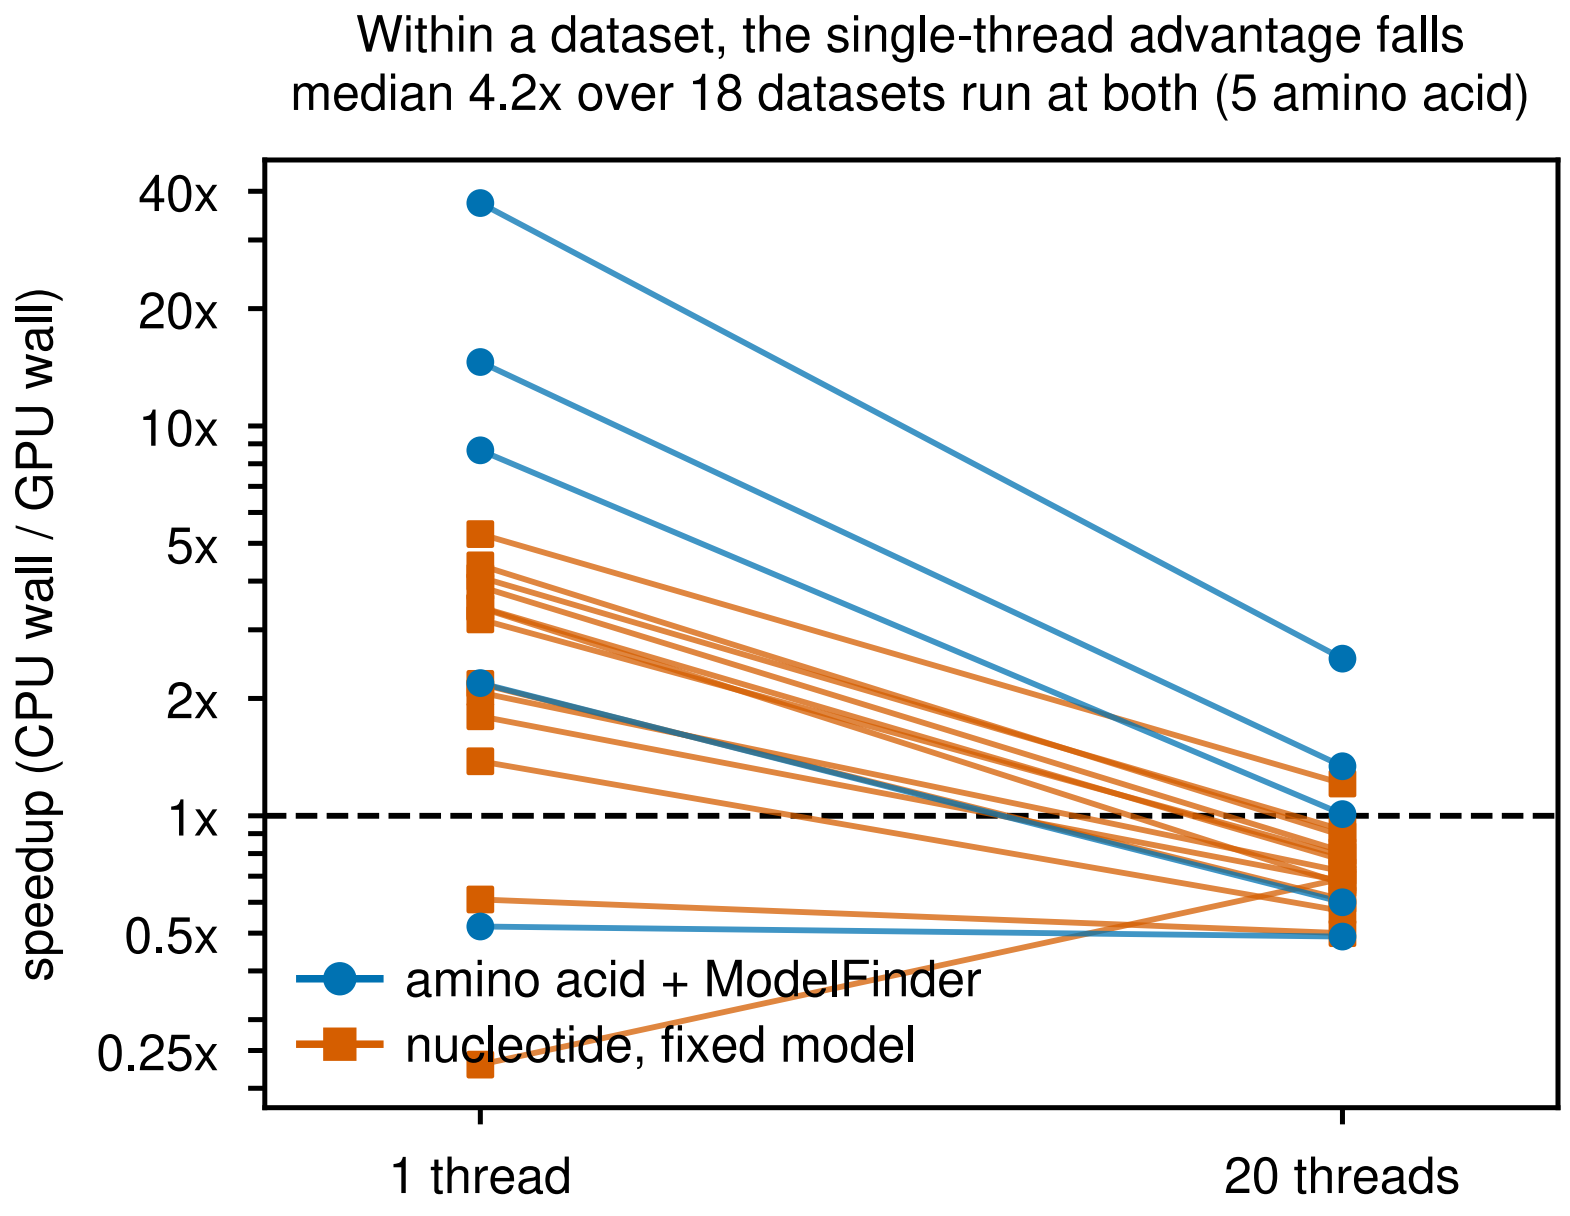
Within a dataset, the single-thread advantage falls
median 4.2x over 18 datasets run at both (5 amino acid)
speedup (CPU wall / GPU wall)
40x
20x
10x
5x
2x
1x
0.5x
0.25x
amino acid + ModelFinder
nucleotide, fixed model
1 thread
20 threads